



# Recipe of Gradient Descent

*Make **Grey** Out of **Black***

(Nov 16, 2016)

YANG Jiancheng



# Outline

- **I. Notes on Gradient Descent**
- **II. Recipe on SGD Family**
  - **Overview**
  - **Momentum**
  - **Nesterov accelerated gradient (NAG)**
  - **Separate / Adaptive Learning Rates: Adagrad, RMSprop, Adadelata and Adam**
- **III. Second Order Method**
- **IV. A Bag of Tricks**



# • I. Notes on Gradient Descent

- Brief Review

## a) Full Batch



$$\theta = \theta - \eta \cdot \nabla_{\theta} J(\theta)$$

## b) SGD (Online)

$$\theta = \theta - \eta \cdot \nabla_{\theta} J(\theta; x^{(i)}; y^{(i)})$$

## c) Mini Batch SGD

$$\theta = \theta - \eta \cdot \nabla_{\theta} J(\theta; x^{(i:i+n)}; y^{(i:i+n)})$$



Linear Neuron Error Surface



# • I. Notes on Gradient Descent

- Highlights

- a) Key Idea behind Stochastic / Mini Batch

- ☐ *Average the moving*
- ☐ Algorithmically: not only on the batch size!

- b) Gradient check ([More on Gradient checks](#))

- ☐ Use the centered formula  $\frac{df(x)}{dx} = \frac{f(x+h) - f(x-h)}{2h}$
- ☐ Compare the relative error
- ☐ Check double float and numerical issue
- ☐ Kinks

- c) Sanity check on the implement:



- ☐ Assume the initial loss
- ☐ Overfit a tiny dataset



# • I. Notes on Gradient Descent

- Issue

- a) Picky learning rate
- b) For global learning rate, hard to:
  - *move faster* when the distance is long, while the gradient is *small*
  - *slow down* when the distance is short, while the gradient is *large*
- c) Saddle points


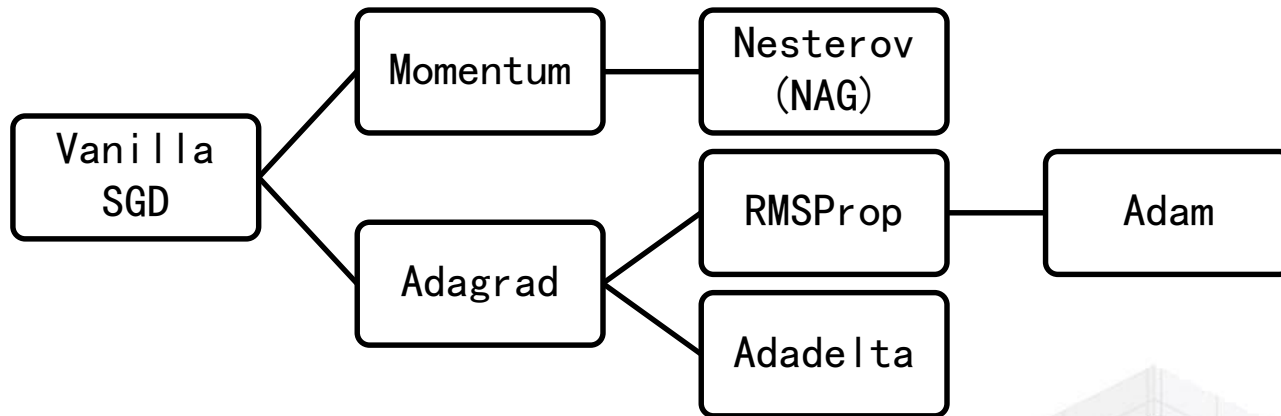


For non-linear neuron, the surface is locally quadratic, with same speed issue.




## • II. Recipe on SGD Family

- Overview



Optimization Speed (contours view)



Optimization Speed (saddle point)




## • II. Recipe on SGD Family

- Momentum



SGD without momentum



SGD with momentum


$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta)$$

$$\theta = \theta - v_t$$




## • II. Recipe on SGD Family

- Nesterov accelerated gradient (NAG)



$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta)$$



$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta - \gamma v_{t-1})$$

$$\theta = \theta - v_t$$

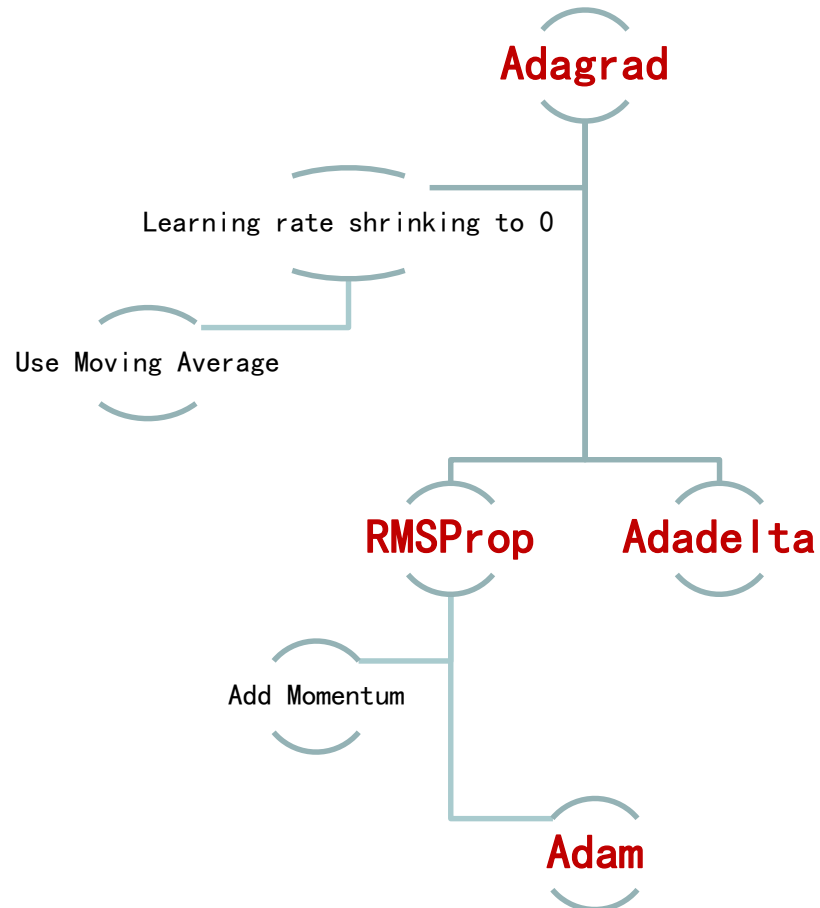




## • II. Recipe on SGD Family

- Separate / Adaptive Learning Rates

**Key idea:** Performing larger updates for infrequent and smaller updates for frequent parameters





## • II. Recipe on SGD Family

- Separate / Adaptive Learning Rates: Adagrad

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{G_t + \varepsilon}}$$

$$G_t = \sum_{\tau=0}^{t-1} g_{\tau}^2$$

- Too aggressive and stops learning too early
- Accumulation of the squared gradients increased  $G_t$  monotonically
- Learning rate shrinks to infinitesimally small



## • II. Recipe on SGD Family

- Separate / Adaptive Learning Rates: RMSProp

$$E[g^2]_t = 0.9E[g^2]_{t-1} + 0.1g_t^2$$

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{E[g^2]_t + \epsilon}} g_t$$

- Proposed by Geoff Hinton in his Coursera class
- Use the decay to simulate the moving windows
- Good choice for Recurrent Neural Network



## • II. Recipe on SGD Family

- Separate / Adaptive Learning Rates: Adadelta

$$RMS[g]_t = \sqrt{E[g^2]_t + \epsilon} \quad E[g^2]_t = \gamma E[g^2]_{t-1} + (1 - \gamma)g_t^2$$

$$RMS[\Delta\theta]_t = \sqrt{E[\Delta\theta^2]_t + \epsilon} \quad E[\Delta\theta^2]_t = \gamma E[\Delta\theta^2]_{t-1} + (1 - \gamma)\Delta\theta_t^2$$

$$\Delta\theta_t = -\frac{RMS[\Delta\theta]_{t-1}}{RMS[g]_t}g_t$$

$$\theta_{t+1} = \theta_t + \Delta\theta_t$$

- No need for a initial learning rate
- Look like RMSProp, developed separately



## • II. Recipe on SGD Family

- Separate / Adaptive Learning Rates: Adam

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t} \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t$$

- Added momentum to gradient
- Use bias correction mechanism



## • III. Second Order Method

### a) Newton's method

$$x \leftarrow x - [H f(x)]^{-1} \nabla f(x)$$

- Hard to get the Hessian

### b) L-BFGS

- Computed over the entire training set


### c) **In practice**, it is currently not common to see L-BFGS or similar second-order methods applied to large-scale Deep Learning



## • IV. A Bag of Tricks

### a) Turn down the learning rate

- A quick win, but slower learning
- Don't too often!



### b) Normalization and Batch Normalization

### c) Modest Initialization

### d) Prefer random search to grid search

### e) Shuffling and Curriculum Learning

- Often good to shuffle the training data before every epoch
- Supplying the training data in a meaningful order may actually help

### f) *Early stopping is beautiful free lunch* (Geoff Hinton)

### g) Gradient Noise

- More robust to poor initialization

### h) Model Ensemble



## • IV. A Bag of Tricks

- Code Tips when Training using Keras

- a) `model.fit(shuffle=True)` will shuffle the training data, before every epoch
- b) Default with shuffle
- c) **Order between “validation\_split” and “shuffle”:**
  - **Will do split first!**
  - If you want to make the validation set different every time, you should keep your own validation set, with “validation\_data”





# Bibliography

- [How to make the learning go faster by Geoffrey Hinton \(Neural Network for Machine Learning Week 6\)](#)
- [An overview of gradient descent optimization algorithms](#)
- [CS231n Course Notes on gradient based optimization](#)



**Thanks for listening!**

